

Characterizing the Performance of Executing Many-tasks on Summit

Matteo Turilli*
Rutgers University
matteo.turilli@rutgers.com

Andre Merzky*
Rutgers University
andre@merzky.net

Thomas Naughton†
Oak Ridge National Laboratory
naughtont@ornl.gov

Wael Elwasif†
Oak Ridge National Laboratory
elwasifwr@ornl.gov

Shantenu Jha
Rutgers University
Brookhaven National Laboratory
shantenu.jha@rutgers.edu

ABSTRACT

Many scientific workloads are comprised of many tasks, where each task is an independent simulation or analysis of data. The execution of millions of tasks on heterogeneous HPC platforms requires scalable dynamic resource management and multi-level scheduling. RADICAL-Pilot (RP) – an implementation of the Pilot abstraction, addresses these challenges and serves as an effective runtime system to execute workloads comprised of many tasks. In this paper, we characterize the performance of executing many tasks using RP when interfaced with JSM and PRRTE on Summit: RP is responsible for resource management and task scheduling on acquired resource; JSM or PRRTE enact the placement of launching of scheduled tasks. Our experiments provide lower bounds on the performance of RP when integrated with JSM and PRRTE. Specifically, for workloads comprised of homogeneous single-core, 15 minutes-long tasks we find that: PRRTE scales better than JSM for $> O(1000)$ tasks; PRRTE overheads are negligible; and PRRTE supports optimizations that lower the impact of overheads and enable resource utilization of 63% when executing $O(16K)$, 1-core tasks over 404 compute nodes.

ACM Reference Format:

Matteo Turilli, Andre Merzky, Thomas Naughton, Wael Elwasif, and Shantenu Jha. 2019. Characterizing the Performance of Executing Many-tasks on Summit. In *Proceedings of Third Annual Workshop on Emerging Parallel and Distributed Runtime Systems and Middleware (IPDRM 2019)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

*Co-first author

†This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, worldwide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (<http://energy.gov/downloads/doe-public-access-plan>).

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

IPDRM 2019, Nov. 22nd, 2019, Denver, Colorado, USA,

© 2019 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Advances in high-performance computing (HPC) have traditionally focused on scale, performance and optimization of applications with a large but single task. Workloads of many scientific applications however, are comprised of many tasks – where each task is an independent computing unit running on one or more nodes, which must be collectively executed and analyzed.

As an example, *ensemble-based* computational methods have been developed to address limitations in single molecular dynamics simulations, where parallelization is limited to speeding up each individual, serialized, time step[12]. Unlike traditional high-throughput “embarrassingly” parallel workloads, these workloads require modest task coordination and inter-task communication (though very infrequent relative to intra-task communication). Multiple simulation tasks are executed concurrently, and various physical or statistical principles are used to combine the output of these tasks, often iteratively and asynchronously. Tens to hundreds of thousands of such tasks are currently needed to adequately sample or investigate the physical phenomenon of interest. Proper sensitivity analysis and uncertainty quantification can increase the total number of tasks by several orders of magnitude.

The execution of millions of tasks on modern HPC platforms is faced with many challenges. A tension exists between the workload’s resource utilization requirements and the capabilities of traditional HPC system software. It requires flexible and dynamic resource management of heterogeneous many-core nodes.

The Pilot abstraction decouples workload specification, resource management, and task execution. Pilot systems – implementations of the Pilot abstraction, submit job placeholders (i.e. pilots) to the resource scheduler. Once active, the pilot accepts and executes tasks directly submitted to it. Tasks are thus executed within time and space boundaries set by the resource scheduler. By implementing multi-level scheduling and late-binding, Pilot systems lower task scheduling overhead, enable higher task execution throughput, and allow greater control over the resources acquired to execute workloads. The pilot must interact with and is dependent on system software to manage the task execution.

RADICAL-Pilot (RP) is a Pilot system that implements the pilot paradigm as outlined in Ref. [18, 26]. RP is implemented in Python and provides a well defined API and usage modes, and is being used by applications drawn from diverse domains, from earth sciences

and biomolecular sciences to high-energy physics. RP directly supports their use of supercomputers or is used as a runtime system by third party workflow or workload management systems [25].

In this paper, we characterize the performance of executing many tasks using RP when it is interfaced with JSM and PRTTE on Summit – a DOE leadership class machine and currently the top ranked supercomputer on the Top 500 list. Summit has 4,608 nodes IBM POWER9 processors and each node has 6 NVIDIA Volta V100s, with a theoretical peak performance of approximately 200 petaFLOPS. JSM is part of LSF and provides services for starting tasks on compute resources; PRTTE provides the server-side capabilities for a reference implementation of the process management interface for ExaScale (PMIx). Specifically, we describe and investigate the baseline performance of the integration of RP [18] with JSM and PRRTE. We experimentally characterize the task execution rates, various overheads, and resource utilization rates.

2 BACKGROUND

We characterize the performance of integrating RADICAL-Pilot (RP) and PMIx Reference RunTime Environment (PRRTE), and RP and IBM Job Step Manager (JSM). These enable the concurrent execution of thousands of application tasks on Summit.

2.1 Process Management Interface for Exascale

The Process Management Interface for Exascale (PMIx) [4] is an open source standard that extends the prior PMI v1 & v2 interfaces used to launch tasks on compute resources. PMIx provides a method for tools and applications to interact with system-level resource managers and process launch mechanisms. PMIx provides a bridge between such clients and underlying execution services, e.g., process launch, signaling, event notification. The clients communicate with PMIx enabled servers, which may support different versions of the standard. PMIx can also be used as a coordination and resource discovery mechanism, e.g., machine topology information. An implementation of the PMIx standard is provided by the OpenPMIx project [3] as a software library that contains the programming interfaces needed to use the standard. The OpenPMIx project also provides a reference implementation of a PMIx enabled runtime (described in Section 2.3).

2.2 IBM Job Step Manager

The IBM Spectrum Load Sharing Facility (LSF) software stack is used to manage the resources for the Summit system. This includes a job scheduler that manages resources and provides allocations based on user provided submissions. The Job Step Manager (JSM) provides services for starting tasks on compute resources within an allocation [2]. The `jsrun` command enables a user to launch an executable on the remote nodes within a user’s job allocation.

When the user is allocated a collection of compute nodes by the batch system, a daemon (`jsmd`) is launched on each of the compute nodes in the allocations. These daemons are then responsible for launching the user’s processes on their respective nodes in response to future `jsrun` commands. There are two startup modes for launching the `jsmd` daemons: *SSH mode* and *non-SSH mode* [2]. As the name suggests, when running in *SSH mode*, Secure Shell is used for launching the `jsmd` processes on the remote nodes of

the allocation. The other option leverages the IBM Cluster Systems Manager (CSM) infrastructure to bootstrap the JSM daemons within the allocation. Currently, the default mode on Summit is to use CSM. Once the JSM daemons are launched, internally the daemons use PMIx [4] to launch, signal, and manage processes on the compute resources.

2.3 PMIx Reference RunTime Environment

A reference implementation of the PMIx server-side capabilities is available via the PMIx Reference RunTime Environment (PRRTE) [11]. PRRTE leverages the modular component architecture (MCA) that was developed for Open MPI [15], which enables execution time customization of the runtime capabilities. The PRRTE implementation provides a portable runtime layer that users can leverage to launch a PMIx server.

PRRTE includes a persistent mode called Distributed Virtual Machine (DVM), which uses system-native launch mechanisms to bootstrap an overlay runtime environment that can then be used to launch tasks via the PMIx interface. This removes the need to bootstrap the runtime layer on each invocation for task launch. Instead, after the launch of the DVM, a tool connects and sends a request to start a task. The task is processed and then generates a launch message that is sent to the PRRTE daemons. These daemons then launch the task. Internally, this task tracking is referred to as a *PRRTE job*, not to be confused with the batch job managed by the system-wide scheduler. The stages of each PRRTE job are tracked from initialization through completion.

We can roughly divide the lifetime of a PRRTE job into the following stages (marked by internal PRRTE state change events): (i) `init_complete` to `pending_app_launch`—time to setup the task and prepare launch details; (ii) `sending_launch_msg` to `running`—time to send the process launch request to PRRTE daemons and to enact them on the target nodes; and (iii) `running` to `notify_complete`—duration of the application plus time to collect task completion notification.

In our experiments, we record the time for the transition between these stages to provide insights on the time spent in the runtime layer when running tasks driven by RP. It should be noted that these phases do not include time between the user launching a PRRTE task and PRRTE initiating processing for this task (e.g., due to file system delays or dynamic libraries loading).

2.4 RADICAL-Pilot

RP [18] is a runtime system designed to decouple resource acquisition from task execution. As every pilot system, RP acquires resources by submitting a batch job, then bootstraps dedicated software components on those resources to schedule, place and launch application tasks, independent from the machine batch system [26]. Scheduling, placing and launching capabilities are specific to each HPC platform, which makes supporting diverse platforms with the same code base challenging. RP can execute single or multi core/GPU tasks within a single compute node, or across multiple nodes. RP isolates the execution of each tasks into a dedicated process, enabling concurrent and sequential execution of heterogeneous tasks by design.

RP is a distributed system designed to instantiate its components across available resources, depending on the platform specifics.

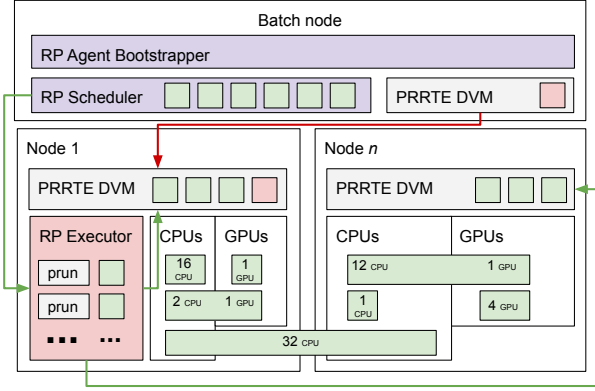


Figure 1: Deployment of RP on Summit with PRRTE/DVM.

Each components can be individually configured so as to enable further tailoring while minimizing code refactoring. RP uses RADICAL-SAGA [19] to support all the major batch systems, including Slurm, PBSPro, Torque and LSF. RP also supports many methods to perform node and core/GPU placement, process pinning and task launching like, for example, aprun, JSM, PRRTE, mpirun, mpiexec and ssh.

RP is composed of two main components: Client and Agent. Client executes on any machine while Agent bootstraps on one of Summit’s batch nodes. Agent is launched by a batch job submitted to Summit’s LSF batch system via RADICAL-SAGA. After bootstrapping, Agent pulls bundles of tasks from Client, manages the tasks’ data dependences if any, and then schedules tasks for execution via either JSM/LSF or PRRTE/DVM.

How Agent deploys on Summit depends on several configurable parameters like, for example, number of sub-agents, number of schedulers and executors per sub-agent, method of communication between agent and sub-agents, and method of placing and launching tasks for each executor of each sub-agent. A default deployment of Agent instantiates a single sub-agent, scheduler and executor on a batch node of Summit. The executor calls one `jsrun` command for each task, and each `jsrun` uses the JSMD demon to place and launch the task on work nodes resources (thread, core, GPU).

Fig. 1 shows the deployment of RP/PRRTE Agent on a batch node and one sub-agent on a compute node. In this configuration, RP uses SSH to launch sub-agents on compute nodes and then PRRTE/DVM to place and launch tasks across compute nodes. This configuration enables the sub-agent to use more resources and, as shown in the next section, improves scalability and performance of task execution. Note that, independent from the configuration and methods used, RP can concurrently place and launch different types of tasks that use different amount and types of resources.

3 PERFORMANCE CHARACTERIZATION

The performance space of RP is vast, including the execution of both homogeneous and heterogeneous tasks, with and without data dependences and at both small and large scales. We thus divide our performance characterization in three phases: (1) scaling the concurrent execution of short, single-core tasks with both JSM and PRRTE; (2) optimizing baseline performance for homogeneous real-life workloads with the best performing between JSM and PRRTE; (3) tailoring performance to data-intensive, compute-intensive and

GPU-intensive workloads. We present the results of the first phase, offering a baseline that we use to drive our development process.

Task here indicates a self-contained executable, executed as one or more processes on the operating system of a Summit compute node. RP, JSM and PRRTE introduce time overheads in tasks execution. These systems require time to schedule, place and launch the task executions. We quantify and compare these overheads, measuring how they change with the scaling of the number of concurrently executed tasks and the amount of used resources.

We differentiate between individual overheads and the integration of these overheads over the execution of the workload. Individual overheads account for the amount of time that single operations add to the execution time of a task. For example, how much time RP takes to schedule a task or PRRTE takes communicating to launch that task. Aggregated overheads indicate how much time performing a group of operations adds to the execution of all the workload tasks. Aggregated overheads account for the overlapping of multiple concurrent operations. For example, given 10 tasks, a scheduling rate of 1 task/s and a scheduling time of 5s for task, the aggregated scheduling overhead would be 15s for full concurrency, while the individual scheduling overhead for each task would be 5s.

The aggregation of the individual overheads across the entire execution determines how available resources are utilized when executing the workload. RP, JSM and PRRTE require resources to perform their operations and some of these operations may partially or globally block the use of available resources. We measure resource utilization showing the portion of resources used or blocked by each system and the percentage of available resources utilized to execute the workload.

3.1 Experiments Design

We perform 4 experiments to measure and compare the performance of RP, JSM and PRRTE when concurrently executing many-tasks workloads on Summit. Task execution requires assigning suitable resources to the tasks, placing them on resources (i.e., a specific compute node, core, GPU or hardware thread) and then launching the execution of those tasks. RP tracks both tasks and available resources, scheduling the former onto the latter; JSM or PRRTE enact task placement and launching.

Experiment 1 quantifies the aggregated overhead of RP measured as the time required to acquire the workload and scheduling its tasks on the available resources. Experiment 1 measures this overhead with both JSM and PRRTE. Experiment 2 quantifies the aggregated overhead of JSM and PRRTE measured as the time from when they receive the tasks from RP to when the tasks start to execute. Experiment 3 measures RP and PRRTE aggregated overheads and resource utilization for scales beyond those currently supported by JSM. Experiment 4 shows the performance improvement obtained by reducing the overheads measured in Experiments 1–3.

In Experiments 1–4 we execute workloads with between 2 and 16384 tasks, each requiring 1 core and executing for 900s. Given Summit architecture, we utilize between 1 and 410 compute nodes, i.e., 42 to 17220 cores, using the SMT1 setting for Summit nodes. Our experiments maximize execution concurrency and therefore the pressure on RP, JSM and PRRTE capabilities. Any lower degree of execution concurrency would require less scheduling, placements and executions, resulting in a better performance of our systems. As




Figure 2: Measured and ideal total execution time (TTX) of the workloads of Experiments 1–3 (green) and 4 (gray).

such, our experiments measure the baseline of the combined scaling performance of RP, JSM or PRRTE on Summit for homogeneous compute-intensive, multi-tasks workloads.

Experiments 1–4 make a set of reasonable simplifications: each task executes the stress command for exactly 900s, a trade off between core allocation cost and the need to stress RP, JSM and PRRTE performance. We do not perform I/O operations as they would be irrelevant to our characterization. JSM and PRRTE do not manage task-level data while RP only links data on the available filesystems and ensures locality of output data. In this context, data performance depends on the storage systems and the executable capabilities and should be independently characterized.

Analogously, in our experiment we do not use real workloads executables. RP, JSM and PRRTE make sure that the executable of a task is launched on the required resources but play no role on their ongoing execution. The executable is self-contained and completely independent from RP, JSM and PRRTE. Thus, the measurements we present apply to every homogeneous workload, independent of the scientific domain, and the specifics of the code executed.

For our experiments we use JSM/jsrun 10.03.00, built for OpenPMIx 3.1.3rc1; PRRTE v1.0.0 with 2 minor patches, built for OpenPMIx 3.1.3; and RADICAL-Cybertools v0.70.3. The data, analysis and code of our experiments is available at [1].

Fig. 2 shows the total execution time (TTX) of the workloads of Experiments 1–4. The black line indicates the ideal execution time, when both software and hardware overheads are assumed to be zero. As expected, the aggregated overheads of the execution increase with the number of tasks and compute nodes utilized. The last column shows Experiment 4 and the marked improvement made by addressing the overheads measured in Experiments 1–3.

3.2 Experiment 1: RP Aggregated Overhead

Fig. 3 shows the RP aggregated overhead when using either JSM or PRRTE to place and launch between 1 and 1024 single-core tasks on between 1 and 49 Summit compute nodes.




Figure 3: Experiment 1: RP aggregated overheads when scheduling 2–1024 single-core tasks on 1–49 compute nodes on Summit.

The mean of the aggregated overhead of RP grows exponentially with scale but differences can be noted between JSM and PRRTE when executing 2 to 8 tasks. With JSM, the aggregated overhead is relatively stable but with large variability at 2 and 32 tasks. With PRRTE, the aggregated overhead grows with minimal variability. Assuming ideal concurrency and resource availability, all the tasks would concurrently execute in 900 seconds. Comparatively, the mean aggregated overhead of RP is < 5% of the ideal total execution time (TTX) with JSM, and < 25% with PRRTE.

Across all scales, the mean of the aggregated overhead of RP is consistently higher with PRRTE than with JSM. This is due to a communication delay we introduced when pushing tasks to PRRTE. RP task scheduling rate is higher than PRRTE task ingestion rate and exceeding it causes PRRTE to fail. We thus slow down RP scheduling rate by introducing an artificial 0.1s delay per task.

PRRTE Wait in Fig. 3 shows the portion of RP aggregated overhead which is due to the delay we introduced. As we are measuring an aggregated overhead, the delays accumulates across the whole execution. PRRTE Wait dominates the RP overhead, showing how, in relative terms, PRRTE overhead is smaller than the one of JSM. Accounting for PRRTE Wait, the mean aggregated overhead of RP is below 3% of the ideal execution time with PRRTE.

We used test runs to empirically derive the amount of delay to introduce in RP when communicating with PRRTE. Exceeding the submission rate with PRRTE leads to tasks submission errors that RP could recover at the cost of reduced utilization. At a rate of 10 tasks/second we observe stable operation of the PRRTE DVM.

Similar test runs uncovered the failure modes of JSM. Originally, exceeding the sustainable rate of calls to JSM caused the LSF daemon to unrecoverably fail. This crashed the LSF job with which RP acquired its resources, causing the failure of the workload execution. Recent updates to LSF on Summit resolved those issues, and no delays are required for sequential JSM invocations.

3.3 Experiment 2: JSM and PRRTE Aggregated Overheads

Fig. 4 shows the aggregated overheads of JSM and PRRTE when executing the same workloads as Experiment 1. Starting from 4/1 tasks/node, JSM has a smaller aggregated overhead compared to PRRTE. PRRTE aggregated overhead grows exponentially with the number of tasks and nodes while JSM shows a less well-defined trend across scales. This is because the delay introduced in RP makes the aggregation of PRRTE individual overheads additive: the delay is longer than PRRTE task placement and launch time.

Characterizing the Performance of Executing Many-tasks on Summit




Figure 4: Experiment 2: JSM and PR RTE aggregated overheads when placing and launching 1–2048 single-core tasks on 1–49 nodes.




Figure 5: Experiment 3: Aggregated overhead of RP and of PR RTE when executing 1024–16384 single-core tasks on 26–410 nodes.

The aggregated overheads of RP, JSM and PR RTE do not sum to the total overhead as measured in Fig.2. RP can schedule tasks while other tasks are been assigned and launched by JSM or PR RTE. Thus, the aggregated overhead of RP and those of JSM or PR RTE may overlap during execution, resulting as a single overhead when projected on TTX.

Both aggregated overheads plateau below 1024 tasks and 25 nodes. This is due to task failure: both JSM and PR RTE can execute up to 967 tasks. Above that, the remaining tasks fail, creating a steady overhead. This upper bound depends on the limit to 4096 open files imposed by the system on a batch node. This results in a maximum of $O(1000)$ tasks as each task consumes at least three file descriptors: standard input, output, and error files.

3.4 Experiment 3: RP and PR RTE at Scale

We overcome the limit of $O(1000)$ concurrent tasks by running multiple instances of RP executor onto compute nodes and reconfiguring the open files limit to 65536. This allows up to ~22000 concurrent tasks per executor but this approach works only with PR RTE. The open files limits cannot be increased with JSM, and JSM becomes unstable with concurrent RP executors. Thus, we could not execute > 967 concurrent task with JSM.

Fig. 5 shows the behavior of the aggregated overheads of both RP and PR RTE at scale. As already observed in Experiment 2, these overheads grow exponentially and, for RP, the waiting time introduced when communicating with PR RTE remains dominant.

16384/401 (single-core) tasks/nodes is the limit currently supported by RP/PR RTE integration. At 32768/815 tasks/nodes, the DVM of PR RTE crashes, likely due to the excessive number of communication channels concurrently opened. Fig. 2 (bottom) shows that the combination of RP and PR RTE aggregated overhead becomes dominant over TTX at 8192/202 tasks/nodes so scaling beyond 16384/410 tasks/cores would not be effective.

The aggregated overheads of RP and PR RTE, alongside the TTX of the workload execution are stable. The variance across runs is minimal, indicating consistent executions across time. Further, experiments 1–4 executed ~200000 tasks without failure, a measure of the robustness of the integration between RP and PR RTE and, up to 967 concurrent tasks, of RP and JSM.

Different from JSM and LSF, PR RTE is an open source project. This allows us to profile the execution of each task inside PR RTE’s DVM. We collect 40 timestamps, that can be grouped pairwise to provide up to 20 sequential durations for each task execution. Profiles allow to isolate overheads both at individual and aggregated levels, enabling to separate in Fig. 5 (bottom) the aggregated task execution overhead and that of the DVM.

Fig. 7 (top) shows the breakdown of the aggregated overhead of PR RTE’s DVM for the execution of Experiment 3 workloads. The dominant aggregated overhead of PR RTE is the time taken to communicate to the demons on each compute node that a task is ready to execute. As already noted, this overhead is the sum of each individual overhead as the rate at which the tasks are queued for execution by RP to PR RTE is too slow to create any overlapping among communication of initiating the execution of tasks.

Fig. 7 (bottom) shows the time taken to communicate the execution of each task within PR RTE’s DVM. The average time taken by each individual overhead is 0.034s, and a standard deviation of 0.047s. The outliers are likely produced by an accumulation in the communication buffer but over the 16384 tasks, the time taken by this communication is mostly stable around the mean. For 16384 tasks, the individual overheads sum up to 570s, ~17% of the TTX of the workload (as shown in Fig. 2 (bottom)).

3.5 Resource Utilization

We measure for how much time each available computing resource is used, blocked or left idling during the execution of a workload. We focus only on the runs with RP/PR RTE as relevant behavior is measured only at the largest scales of our experiments.

Resources become available as soon as Summit’s scheduler bootstraps the job submitted to Summit’s batch system on one of the requested compute nodes. From then on, we distinguish among three main consumers of those resources: RP, PR RTE and the workload. Each consumer can use the resources to perform computations, block the resources while a computation is performed, or resources can idle because they have no active consumers.

The percentage of resource consumed indicate how much of the resources has been actually used to execute the scientific payload and, in case, where resources have been wasted while being blocked by other operations. It is the most relevant measure of the baseline performance of RP and PR RTE integration.

Fig. 6 shows resource utilization (RU) for Experiment 3. *Agent Nodes* (dark orange) indicates the resources allocated exclusively to RP. *Pilot Startup* (blue) shows the time in which the resources are blocked while RP starts up; and *Warmup* (light orange) the time in which resources are blocked by RP while collecting tasks and scheduling them for execution. *Prepare Execution* (light green) indicates the resources blocked while waiting to communicate with PR RTE for task execution; *Execution Cmd* (black) marks the time in which tasks use resources for execution. *Draining* (dark green) indicates the resources blocked while waiting to drain tasks upon their




Figure 6: Experiment 3: Resource utilization as the time in which each available core has been utilized or blocked by RP, PRRTE or the workload execution.




Figure 7: Experiment 3: Dominant aggregated overheads of PRRTE when executing 1024/26–16384/410 tasks/nodes on Summit.

completion; and *Pilot Termination* (light gray) shows the resources blocked while terminating the pilot.

Consistent with the overhead analysis, execution preparation (light green) and execution draining (dark green) progressively dominate RU with scale. Execution preparation corresponds to the wait time introduced in RP and draining time is specular to wait time: the slower is the rate at which tasks are started, the slower is the rate at which they can be drained. Both starting and draining time are blocking operations that, at the current supported launching rates, result in large amount of available resources not to be used for the execution of the workload.

Execution RP, *Execution PRRTE* and *Unschedule* are too small to be seen when compared to the other measures of RU. This indicates that PRRTE has no appreciable impact over RU during the workload execution. RP impact is noticeable for exclusive resource allocation (a whole compute node), and for blocking available resources while bootstrapping and preparing the execution. *Pilot termination* is visible only at the lower scales as it has a mostly constant impact on RU. Fig. 6 shows several horizontal green lines, cutting across each plot, indicating resources idling across the whole execution. In Experiment 3, these resources are GPUs as our workload does not use them.

Table 1 details our measures of RU as percentage of the available resources. Resources used by RP are independent of resource size, thus the percentage of resource utilized by RP decreases with the size of the pilot. Similarly, the percentage of available resource blocked while starting the pilot decreases with scale as startup time is relatively invariant across pilot sizes. The resources blocked while

“warming up”, i.e., collecting and scheduling tasks, significantly increases from 8192 tasks onwards. This is mainly due to scheduling efficiency as RP scheduler performance depends on the amount of available resources. Consistently with what observed in Figs. 7 and 6, PRRTE has a negligible impact on RU across all the scales.

3.6 Discussion

Experiments 1–3 and the metrics time to execution (TTX) and resource utilization (RU) offer three main insights: (1) performance baseline of RP with JSM or PRRTE on Summit up to the scale currently supported; (2) characterization of aggregated and individual overheads of RP, JSM and PRRTE; and (3) performance evaluation of RP for TTX and RU.

One of the main goals of our performance baseline is to guide the engineering with which we integrate RP, JSM and PRRTE. The analysis we presented shows that the waiting time between RP and PRRTE communication is the main barrier to scalability. Fig. 3 and Fig. 4 show that PRRTE Wait is the dominant aggregated overhead. The analysis of Fig. 7 showed how this waiting time determines the aggregation of PRRTE overheads into the sum of non-overlapping task overheads. Further, Fig. 6, shows that the waiting time reduces up to 3/4 the resources that the workload can utilize for execution.


The delay we introduced is conservative so to guarantee no failure in task execution. In many real life use cases, task failure can be managed with fault tolerance as done, for example, with RADICAL Ensemble Toolkit (EnTK) and RP when executing seismic inversion on Titan [7]. There, we scaled the execution up to 131,000 cores, resubmitting ~15% of tasks due to diverse types of failure. We used the same approach on Summit: eliminating RP waiting time with PRRTE, lead to between 3 and 10% task failure rate with 2 total execution failures out of 8 runs.

Based on the analysis of the failures we recorded, we configured PRRTE to use a flat communication hierarchy and ssh as its communication channel. This reduced the internal performance of PRRTE and limited the total amount of concurrent tasks that it can handle to ~20000 but it also allowed a more aggressive communication rate between RP and PRRTE. In Experiment 4, we were able to reduce the waiting time from 0.1s to 0.01s and to use 4 concurrent sub-agents for RP. This increased the rate of communication between RP and PRRTE both for each single sub-agent and globally, due to the concurrency among sub-agents.

Fig. 8 shows how this dramatically improved RU. Compared to the same run on Experiment 3, Experiment 4 reduced the mean of TTX from 3236s to 1296s, the mean of aggregated RP overhead from 2648s to 522s, and the mean of aggregated overhead of PRRTE from 2228s to 341s.

Table 1: Experiments 3–4: Resource utilization (RU) expressed as the percentage of resources used or blocked by RP, PR RTE and the workload. Last line shows optimized run of Experiment 4.

Tasks / Nodes	Agent Nodes	Pilot Startup	Warmup	Prep. Execution	Exec. RP	Exec. PR RTE	Exec. Cmd	Unschedule	Draining	Pilot Termination	Idle
1024 / 26	3.846%	3.630%	1.680%	4.510%	0.016%	0.002%	73.999 %	0.001%	6.149%	0.812%	5.355%
2048 / 51	1.961%	3.622%	1.603%	9.800%	0.011%	0.004%	65.313 %	0.000%	11.356%	0.867%	5.462%
4096 / 101	0.990%	2.698%	1.398%	16.178%	0.013%	0.002%	54.797 %	0.000%	17.798%	0.534%	5.593%
8192 / 202	0.495%	2.076%	1.954%	23.375%	0.021%	0.002%	39.990 %	0.001%	25.570%	0.396%	6.120%
16384 / 410	0.244%	1.271%	3.309%	28.779%	0.021%	0.002%	25.596 %	0.001%	32.752%	0.256%	7.771%
16384 / 410	1.013%	3.265%	6.314%	2.345%	2.421%	4.988%	63.557 %	0.286%	11.526%	0.800%	3.485%

**Figure 8: Experiment 4: Resource utilization of a 16384/404 tasks/nodes run with optimized RP/PR RTE integration.**

The last line of Table 1 shows the details of RU for Experiment 4. RU of the workload improved from 25% to 63% while the RU of preparing execution decreased from 29% to 2% and RU of idling resources decreased from 8% to 3%. RU of both RP and PR RTE grows: RP deployment requires more time due to the increased number of components instantiated and the higher rate of scheduling; and the increased rate of task scheduling, placement and execution stress more the capabilities of RP and PR RTE implementations.

Eliminating the delay between RP and PR RTE requires to introduce concurrency between the two systems. We prototyped a version of RP that partitions the available resources and uses a DVM for each partition. Tasks will be scheduled across partitions, adding a minimal overhead for meta-scheduling operations. Each DVM will operate on smaller portions of resources, lowering the number of tasks scheduled by RP. In turn, this will eliminate the need to add a waiting time, further reducing the overheads measured in our baseline and improving both TTX and RU.

The main space for further improvement in RP is scheduling efficiency. RP already support a scheduling framework in which algorithms tailored to the executed workload can be used to optimize performance. Nonetheless, in RP the scheduling algorithms are implemented in Python, leading to an unavoidable ceiling on the amount of optimization we can implement at large scales. Prototypes implemented in C show the near complete elimination of scheduling overheads when executing both heterogeneous and homogeneous workloads. Integration with third-party tools like Flux [5] is another promising approach to solve this problem.

PR RTE overheads for individual tasks are both stable and small. This leaves little space of optimization for the execution of many-tasks workloads as those scheduled by RP. The main improvement for PR RTE would be to increase the number of concurrent tasks that can be managed by the DVM but the importance of such an

improvement will decrease once RP can utilize multiple DVMs within the same pilot.

4 RELATED WORK

Pilot systems like GlideinWMS [23], PanDA [13] and DIRAC [24] are used to implement late binding and multi-level scheduling on a variety of platforms. While these systems have been successfully used on HPC machines [14, 16, 17], including on the former ORNL leadership class machine Titan [20], they are currently not available on Summit and do not support either JSM or PR RTE.

Both PR RTE [11] and JSM [21] rely on PMIx to place and launch processes on Summit’s nodes. Many applications are actively working to directly use PMIx to interface with the system management stack to benefit from portable process and resource management capabilities [27]. While PMIx explicitly supports interfacing with command line tools, there are no other pilot systems using PMIx via JSM or PR RTE. MPICH and Hydra [6] offer capabilities similar to PR RTE but are not supported on Summit.

Pilot systems are not the only way to execute many-task applications on HPC machines. JSM and LSF natively support this capability but, as seen in §3, in their current deployment on Summit they cannot scale beyond 1000 concurrent task executions. Flux [5] is a resource manager that provides users with private schedulers on pools of dedicated resources. This enables the task scheduling capabilities of a pilot system, including RP, but requires to be either adopted as the main job manager of the machine or be deployed as part of a pilot system.

METAQ [8, 10] are a set of shell scripts that forms a “middle layer” between the batch scheduler and the user’s computational job scripts and supports task packing. METAQ requires a separate invocation of mpirun (or equivalent) for each task. METAQ has been superseded by mpi_jm [9] — a python library that is linked to applications. In addition to intelligent backfilling and task packing, mpi_jm allows the executable to be launched based upon an affinity with the hardware.

In Ref.[18, 22] we investigated the performance of RP on ORTE — a precursor to PR RTE. Using ORTE, RP was capable of spawning more than 100 tasks/second and the steady-state execution of up to 16K concurrent tasks. Resource utilization was significant lower than with PR RTE and more sensitive to the number of units and unit duration.

5 CONCLUSIONS

We characterized the performance of the integration between RP, JSM and PR RTE on Summit when executing many-task workloads. Our baseline characterizes aggregated and individual overheads for each system, measuring resource utilization for each available

computing core. Our baseline measures the performance for the worst case scenario in which single-core, 15 minutes-long tasks are executed on up to 410 compute nodes of Summit. Further, based on the insight gained by our characterization, we showed the results of our optimization when executing 16384 1-core, 15 minutes-long tasks on up to 404 compute nodes.

Our experiments shows that on Summit: (1) PRRTE enables better scalability than JSM when executing homogeneous many-task applications at scales larger than 987 concurrent task executions; (2) up to the scale currently supported, PRRTE individual overheads are negligible when compared to other overheads; and (3) PRRTE open source code enables optimizations that lower the impact of aggregated overheads over the execution of the considered workloads. Further, we show that RP can effectively integrate with both JSM and PRRTE, imposing manageable aggregated overheads while offering high degrees of configurability. Specifically, we show that once optimized, at the largest scale supported and for the considered workload the integration between RP and PRRTE imposes an overall aggregated overhead of 35% over the total time to execution of the workload. This enables the utilization of 63% of the available resources to execute the given workload.

The presented performance characterization, its analysis, and the implemented optimizations are the foundation of future work with RP, JSM, PRRTE and Summit. The current scale at which RP/PRRTE operate support the development of three use cases: machine learning driven molecular dynamics simulations; machine learning driven drug discovery protocols; and seismic inversion workflows. RP and PRRTE are posed to support several INCITE and Exascale computing projects, accounting for a significant portion of the available allocation on Summit in the next years. To this end, we will enable RP to partition both available resource and workload execution. As seen in §3.6, this will greatly reduce aggregated overheads and improve resource utilization efficiency. Further work will be needed to optimize RP scheduler when managing workload with both spatial and temporal heterogeneity, i.e., those in which task execution time are drawn from a large distribution. The next step will be to characterize performance at increasingly large scales, while measuring and addressing the bottlenecks for heterogeneous workloads executed both concurrently and sequentially on the same pilot.

6 ACKNOWLEDGMENTS

We would like to thank other members of the PMIx community, and Ralph Castain in particular, for the excellent work that we build upon. This research used resources of the Oak Ridge Leadership Computing Facility at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725. At Rutgers, this work was also supported by NSF “CAREER” ACI-1253644, RADICAL-Cybertools NSF 1440677 and 1931512, and DOE Award DE-SC0016280. We also acknowledge DOE INCITE awards for allocations on Summit.

REFERENCES

- [1] Experiment data and analysis code for this paper, https://github.com/radical-experiments/summit_jsrun_prte. Accessed: 2019-08-31.
- [2] LSF job step manager version 10.3. https://www.ibm.com/support/knowledgecenter/SSWRJV_10.1.0/jsm/10.3/base/jsm_kickoff.html. Accessed: 2019-10-04.
- [3] OpenPMIx project. <https://github.com/openpmix/prte/>. Accessed: 2019-10-04.
- [4] Process management interface for exascale (PMIx) standard. <https://pmix.org/pmix-standard>. Accessed: 2019-10-04.
- [5] AHN, D. H., GARLICK, J., GRONDONA, M., LIPARI, D., SPRINGMEYER, B., AND SCHULZ, M. Flux: A next-generation resource management framework for large hpc centers. In *43rd IEEE International Conference on Parallel Processing Workshops (ICCPW)* (2014), pp. 9–17.
- [6] BALAJI, P., BLAND, W., GROPP, W., LATHAM, R., LU, H., PENA, A. J., RAFFENETTI, K., SEO, S., THAKUR, R., AND ZHANG, J. Mpi user’s guide. *Argonne National Laboratory* (2014).
- [7] BALASUBRAMANIAN, V., TURILLI, M., HU, W., LEFEBVRE, M., LEI, W., MODRAK, R., CERVONE, G., TROMP, J., AND JHA, S. Harnessing the power of many: Extensible toolkit for scalable ensemble applications. In *IEEE International Parallel and Distributed Processing Symposium (IPDPS)* (2018), pp. 536–545.
- [8] BERKOWITZ, E. Metaq: Bundle supercomputing tasks. <https://arxiv.org/abs/1702.06122>.
- [9] BERKOWITZ, E., JANSEN, G., McELVAIN, K., AND WALKER-LOUD, A. Job management with mpi_fm. In *International Conference on High Performance Computing* (2018), Springer, pp. 432–439.
- [10] BERKOWITZ, E., JANSEN, G. R., McELVAIN, K., AND WALKER-LOUD, A. Metaq: Bundle supercomputing tasks. *EPJ Web Conf.*, vol. 175, p. 09007, 2018.
- [11] CASTAIN, R. H., HURSEY, J., BOUTELLIER, A., AND SOIT, D. PMIx: process management for exascale environments. *Parallel Computing* 79 (2018), 9–29.
- [12] CHEATHAM, T. E., AND ROE, D. R. The impact of heterogeneous computing on workflows for biomolecular simulation and analysis. *Computing in Science and Engineering* 17, 2 (2015), 30–39.
- [13] DE, K., KLIMENTOV, A., WENAU, T., MAENO, T., AND NILSSON, P. PanDA: A new paradigm for distributed computing in HEP through the lens of ATLAS and other experiments. Tech. rep., ATL-COM-SOFT-2014-027, 2014.
- [14] FIFIELD, T., CARMONA, A., CASAJÚS, A., GRACIANI, R., AND SEVIOR, M. Integration of cloud, grid and local cluster resources with dirac. In *Journal of Physics: Conference Series* (2011), vol. 331, IOP Publishing, p. 062009.
- [15] GABRIEL, E., FAGG, G. E., BOSILCA, G., ANGSKUN, T., DONGARRA, J. J., SQUYRES, J. M., SAHAY, V., KAMBADUR, P., BARRETT, B., LUMSDAINE, A., CASTAIN, R. H., DANIEL, D. J., GRAHAM, R. L., AND WOODALL, T. S. Open MPI: Goals, concept, and design of a next generation MPI implementation. In *Proceedings, 11th European PVM/MPI Users’ Group Meeting* (Budapest, Hungary, September 2004), pp. 97–104.
- [16] HUFNAGEL, D. Cms use of allocation based hpc resources. In *J. Phys. Conf. Ser.* (2017), vol. 898, p. 092050.
- [17] MAENO, T., AND ET AL. Evolution of the ATLAS PanDA workload management system for exascale computational science. In *Proceedings of the 20th International Conference on Computing in High Energy and Nuclear Physics (CHEP2013)*, *Journal of Physics: Conference Series* (2014), vol. 513(3), IOP Publishing, p. 032062.
- [18] MERZKY, A., TURILLI, M., MALDONADO, M., SANTCROOS, M., AND JHA, S. Using pilot systems to execute many task workloads on supercomputers. In *Workshop on Job Scheduling Strategies for Parallel Processing* (2018), Springer, pp. 61–82.
- [19] MERZKY, A., WEIDNER, O., AND JHA, S. SAGA: A standardized access layer to heterogeneous distributed computing infrastructure. *Software-X* (2015).
- [20] OLEYNIK, D., PANITKIN, S., TURILLI, M., ANGIUS, A., ORAL, S., DE, K., KLIMENTOV, A., WELLS, J. C., AND JHA, S. High-throughput computing on high-performance platforms: A case study. In *2017 IEEE 13th International Conference on e-Science (e-Science)* (2017), IEEE, pp. 295–304.
- [21] QUINTERO, D., GOMEZ GONZALEZ, M., HUSSEIN, A. Y., AND MYKLEBUST, J.-F. *IBM High-Performance Computing Insights with IBM Power System AC922 Clustered Solution*. IBM Redbooks, 2019.
- [22] SANTCROOS, M., CASTAIN, R., MERZKY, A., BETHUNE, I., AND JHA, S. Executing dynamic heterogeneous workloads on blue waters with radical-pilot. *Cray User Group 2016* (2016).
- [23] SFILIGOI, I. glideinWMS—a generic pilot-based workload management system. In *Proceedings of the international conference on computing in high energy and nuclear physics (CHEP2007)*, *Journal of Physics: Conference Series* (2008), vol. 119(6), IOP Publishing, p. 062044.
- [24] TSAREGORODTSEV, A., GARONNE, V., AND STOKES-REES, I. DIRAC: A scalable lightweight architecture for high throughput computing. In *Proceedings of the 5th IEEE/ACM International Workshop on Grid Computing* (2004), pp. 19–25.
- [25] TURILLI, M., BALASUBRAMANIAN, V., MERZKY, A., PARASKEVAKOS, I., AND JHA, S. Middleware building blocks for workflow systems. *Computing in Science & Engineering* (2019).
- [26] TURILLI, M., SANTCROOS, M., AND JHA, S. A comprehensive perspective on pilot-job systems. *ACM Computing Surveys (CSUR)* 51, 2 (2018), 43.
- [27] VALLÉE, G. R., AND BERNHOLDT, D. Improving support of MPI+OpenMP applications. In *Proceedings of the EuroMPI 2018 Conference* (2018).